




Fine-tune loss curves



Gradient norm per step



Parameter shift magnitude (L2 norm of updates)



Gradient cosine similarity (current vs previous)

