




Fine-tune loss curves




Gradient norm per step



Parameter shift magnitude (L2 norm of updates)



Gradient cosine similarity (current vs previous)



Alignment ratio (avg over task1 samples in batch)



Phase 2 eval during training (every 20 steps)

