




### Fine-tune loss curves




### Gradient norm per step



Parameter shift magnitude (L2 norm of updates)



### Gradient cosine similarity (current vs previous)



Phase 2 eval during training (every 20 steps) and final

