




Teacher fine-tune — task 1 (teacher eval SR: 100.0%)




Teacher fine-tune — task 2 (from base) — teacher eval pending




Student BC loss (continuous student across both tasks)



Phase 1 — distillation rollout vs periodic 10-ep eval (task 1)




Phase 2 — task2 rollout + dual 10-ep evals every 5 outer iters




Student BC — each training phase


Student BC (outer iter t1:0)




Student BC (outer iter t1:1)




Student BC (outer iter t1:2)




Student BC (outer iter t1:3)




Student BC (outer iter t1:4)




Student BC (outer iter t1:5)




Student BC (outer iter t1:6)




Student BC (outer iter t1:7)




Student BC (outer iter t1:8)




Student BC (outer iter t1:9)




Student BC (outer iter t1:10)




Student BC (outer iter t1:11)




Student BC (outer iter t1:12)




Student BC (outer iter t1:13)



Student BC (outer iter t1:14)



Student BC (outer iter t1:15)



Student BC (outer iter t1:16)

