




Teacher fine-tune — task 1 (teacher eval SR: 100.0%)




Teacher fine-tune — task 2 (from base) (teacher eval SR: 100.0%)




Student BC loss (continuous student across both tasks)



Phase 1 — distillation rollout vs periodic 10-ep eval (task 1)



Phase 2 — task2 rollout + dual 10-ep evals every 5 outer iters



Student BC — each training phase

